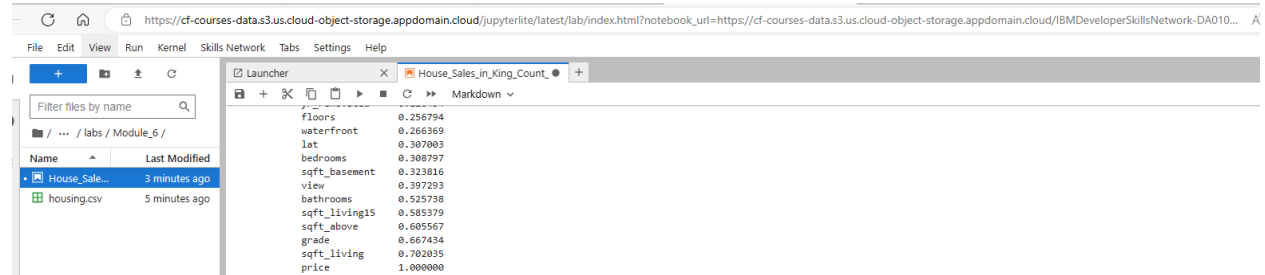


House sales King County – Final Lab Data Analysis for Python – Coursera (HS)

Path:



Question 1

Display the data types of each column using the function `dtypes`. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
[39]: df.dtypes

[39]: Unnamed: 0      int64
      id            int64
      date          object
      price         float64
      bedrooms      float64
      bathrooms     float64
      sqft_living    int64
      sqft_lot       int64
      floors         float64
      waterfront     int64
      view           int64
      condition      int64
      grade          int64
      sqft_above     int64
      sqft_basement  int64
      yr_built       int64
      yr_renovated   int64
      zipcode        int64
      lat            float64
      long           float64
      sqft_living15  int64
      sqft_lot15     int64
      dtype: object

We use the method describe to obtain a statistical summary of the dataframe.
```

Question 2

Drop the columns "id" and "Unnamed: 0" from axis 1 using the method `drop()`, then use the method `describe()` to obtain a statistical summary of the data. Make sure the `inplace` parameter is set to `True`. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
[41]: df.drop("id", axis = 1, inplace = True)
      df.drop("Unnamed: 0", axis = 1, inplace = True)
      df.describe()

[41]:
```

| | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront | view | condition | grade | sqft_above | sqft_basement | yr_built | yr_renovated | zipcode |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|---------------|--------------|--------------|--------------|
| count | 2.161300e+04 | 21600.000000 | 21603.000000 | 21613.000000 | 2.161300e+04 | 21613.000000 | 21613.000000 | 21613.000000 | 21613.000000 | 21613.000000 | 21613.000000 | 21613.000000 | 21613.000000 | 21613.000000 | 21613.000000 |
| mean | 5.400881e+05 | 3.372870 | 2.115736 | 2079.899736 | 1.510697e+04 | 1.494309 | 0.007542 | 0.234303 | 3.409430 | 7.656873 | 1788.390691 | 291.509045 | 1971.005136 | 84.402258 | 98077.939805 |
| std | 3.671272e+05 | 0.926657 | 0.768996 | 918.440897 | 4.142051e+04 | 0.539989 | 0.086517 | 0.766318 | 0.650743 | 1.175459 | 828.090978 | 442.575043 | 29.373411 | 401.679240 | 53.505026 |
| min | 7.500000e+04 | 1.000000 | 0.500000 | 290.000000 | 5.200000e+02 | 1.000000 | 0.000000 | 0.000000 | 1.000000 | 1.000000 | 290.000000 | 0.000000 | 1900.000000 | 0.000000 | 98001.000000 |
| 25% | 3.219500e+05 | 3.000000 | 1.750000 | 1427.000000 | 5.040000e+03 | 1.000000 | 0.000000 | 0.000000 | 3.000000 | 7.000000 | 1190.000000 | 0.000000 | 1951.000000 | 0.000000 | 98033.000000 |
| 50% | 4.500000e+05 | 3.000000 | 2.250000 | 1910.000000 | 7.618000e+03 | 1.500000 | 0.000000 | 0.000000 | 3.000000 | 7.000000 | 1560.000000 | 0.000000 | 1975.000000 | 0.000000 | 98065.000000 |
| 75% | 6.450000e+05 | 4.000000 | 2.500000 | 2550.000000 | 1.068800e+04 | 2.000000 | 0.000000 | 0.000000 | 4.000000 | 8.000000 | 2210.000000 | 560.000000 | 1997.000000 | 0.000000 | 98118.000000 |
| max | 7.700000e+06 | 33.000000 | 8.000000 | 13540.000000 | 1.651359e+06 | 3.500000 | 1.000000 | 4.000000 | 5.000000 | 13.000000 | 9410.000000 | 4820.000000 | 2015.000000 | 2015.000000 | 98199.000000 |

We can see we have missing values for the columns bedrooms and bathrooms

Module 3: Exploratory Data Analysis

Question 3

Use the method `value_counts` to count the number of houses with unique floor values, use the method `.to_frame()` to convert it to a data frame. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
[46]: df['floors'].value_counts().to_frame()
```


```
[46]: floors
1.0    10680
2.0     8241
1.5     1910
3.0      613
2.5      161
3.5         8
```

Question 4

Use the function `boxplot` in the seaborn library to determine whether houses with a waterfront view or without a waterfront view have more price outliers. Take a screenshot of your code and boxplot. You will need to submit the screenshot for the final project.

```
[47]: sns.boxplot(x="waterfront", y="price", data=df)
```

```
[47]: <AxesSubplot:xlabel='waterfront', ylabel='price'>
```




Question 5

Use the function `regplot` in the seaborn library to determine if the feature `sqft_above` is negatively or positively correlated with price. Take a screenshot of your code and scatterplot. You will need to submit the screenshot for the final project.

```
[48]: sns.regplot(x="sqft_above", y="price", data=df, ci=None)
```

```
[48]: <AxesSubplot:xlabel='sqft_above', ylabel='price'>
```



Module 4: Model Development

We can fit a linear regression model using the longitude feature `'long'` and calculate the R^2 .

```
[51]: X = df[['long']]
      Y = df[['price']]
      lm = LinearRegression()
      lm.fit(X,Y)
      lm.score(X, Y)
```

```
[51]: 0.00046769430149007363
```

Question 6

Fit a linear regression model to predict the `'price'` using the feature `'sqft_living'` then calculate the R^2 . Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[52]: X1 = df[['sqft_living']]
      Y1 = df[['price']]
      lm = LinearRegression()
      lm
      lm.fit(X1,Y1)
      lm.score(X1, Y1)
```

```
[52]: 0.4928532179037931
```

Question 7

Fit a linear regression model to predict the `'price'` using the list of features:

```
[53]: features =["floors", "waterFront","lat" ,"bedrooms" ,"sqft_basement" ,"view" ,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]
```

Then calculate the R^2 . Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[54]: X2 = df[features]
      Y2 = df[['price']]
      lm.fit(X2,Y2)
      lm.score(X2,Y2)
```

```
[54]: 0.6576890354915759
```

Question 8

Use the list to create a pipeline object to predict the 'price', fit the object using the features in the list `features` , and calculate the R^2 . Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[56]: pipe=Pipeline(Input)
      pipe
      X = df[features]
      Y = df[['price']]
      pipe.fit(X,Y)
      pipe.score(X,Y)
```

```
[56]: 0.7512051345272872
```

Question 9

Create and fit a Ridge regression object using the training data, set the regularization parameter to 0.1, and calculate the R^2 using the test data. Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[59]: from sklearn.linear_model import Ridge
```

```
[60]: RidgeModel = Ridge(alpha=0.1)
      RidgeModel.fit(x_train, y_train)
      RidgeModel.score(x_test, y_test)
```

```
[60]: 0.647875916393907
```

Question 10

Perform a second order polynomial transform on both the training data and testing data. Create and fit a Ridge regression object using the training data, set the regularisation parameter to 0.1, and calculate the R^2 utilising the test data provided. Take a screenshot of your code and the R^2 . You will need to submit it for the final project.

```
[61]: pr = PolynomialFeatures(degree = 2)
      x_train_pr = pr.fit_transform(x_train[features])
      x_test_pr = pr.fit_transform(x_test[features])

      RidgeModel1 = Ridge(alpha = 0.1)
      RidgeModel1.fit(x_train_pr, y_train)
      RidgeModel1.score(x_test_pr, y_test)
```

```
[61]: 0.7002744263583341
```